

Import the Required libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from keras.models import Model
from keras.layers import LSTM, Activation, Dense, Dropout, Input,
Embedding
from keras.optimizers import RMSprop
from keras.preprocessing.text import Tokenizer
from keras.preprocessing import sequence
from keras.utils import pad_sequences
from keras.utils import to_categorical
from keras.callbacks import EarlyStopping
```

Read the Dataset

```
df = pd.read_csv('/content/spam.csv',delimiter=',',encoding='latin-1')
df.head()
```

	v1	v2	Unnamed: 2
\			
0	ham	Go until jurong point, crazy.. Available only ...	NaN
1	ham	Ok lar... Joking wif u oni...	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN
3	ham	U dun say so early hor... U c already then say...	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN

	Unnamed: 3	Unnamed: 4
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

Preprocessing the Dataset

```
df.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'],axis=1,inplace=True)
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0    v1      5572 non-null    object
 1    v2      5572 non-null    object
dtypes: object(2)
memory usage: 87.2+ KB

X = df.v2
Y = df.v1
le = LabelEncoder()
Y = le.fit_transform(Y)
Y = Y.reshape(-1,1)

X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.15)

max_words = 1000
max_len = 150
tok = Tokenizer(num_words=max_words)
tok.fit_on_texts(X_train)
sequences = tok.texts_to_sequences(X_train)
sequences_matrix = pad_sequences(sequences,maxlen=max_len)

```

Create Model and Add Layers

```

inputs = Input(name='inputs',shape=[max_len])
layer = Embedding(max_words,50,input_length=max_len)(inputs)
layer = LSTM(64)(layer)
layer = Dense(256,name='FC1')(layer)
layer = Activation('relu')(layer)
layer = Dropout(0.5)(layer)
layer = Dense(1,name='out_layer')(layer)
layer = Activation('sigmoid')(layer)
model = Model(inputs=inputs,outputs=layer)

```

```
model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
inputs (InputLayer)	[(None, 150)]	0
embedding (Embedding)	(None, 150, 50)	50000
lstm (LSTM)	(None, 64)	29440
FC1 (Dense)	(None, 256)	16640

activation (Activation)	(None, 256)	0
dropout (Dropout)	(None, 256)	0
out_layer (Dense)	(None, 1)	257
activation_1 (Activation)	(None, 1)	0

```

=====
Total params: 96,337
Trainable params: 96,337
Non-trainable params: 0
=====

```

Compiling the Model

```
model.compile(loss='binary_crossentropy',optimizer=RMSprop(),metrics=['accuracy'])
```

Training the Model

```
model.fit(sequences_matrix,Y_train,batch_size=128,epochs=10,
        validation_split=0.2)
```

```

Epoch 1/10
30/30 [=====] - 11s 361ms/step - loss: 0.0033
- accuracy: 0.9992 - val_loss: 0.1227 - val_accuracy: 0.9884
Epoch 2/10
30/30 [=====] - 8s 273ms/step - loss: 0.0027
- accuracy: 0.9992 - val_loss: 0.1363 - val_accuracy: 0.9884
Epoch 3/10
30/30 [=====] - 8s 277ms/step - loss: 0.0025
- accuracy: 0.9992 - val_loss: 0.1368 - val_accuracy: 0.9905
Epoch 4/10
30/30 [=====] - 8s 272ms/step - loss: 0.0018
- accuracy: 0.9997 - val_loss: 0.1411 - val_accuracy: 0.9895
Epoch 5/10
30/30 [=====] - 8s 272ms/step - loss: 0.0019
- accuracy: 0.9997 - val_loss: 0.1418 - val_accuracy: 0.9895
Epoch 6/10
30/30 [=====] - 8s 271ms/step - loss: 0.9085
- accuracy: 0.9570 - val_loss: 0.2119 - val_accuracy: 0.9863
Epoch 7/10
30/30 [=====] - 8s 272ms/step - loss: 0.0045
- accuracy: 0.9987 - val_loss: 0.1805 - val_accuracy: 0.9905
Epoch 8/10
30/30 [=====] - 8s 271ms/step - loss: 0.0021
- accuracy: 0.9995 - val_loss: 0.1587 - val_accuracy: 0.9905
Epoch 9/10
30/30 [=====] - 8s 270ms/step - loss: 0.0023
- accuracy: 0.9995 - val_loss: 0.1543 - val_accuracy: 0.9895

```

```
Epoch 10/10  
30/30 [=====] - 8s 272ms/step - loss: 0.0027  
- accuracy: 0.9995 - val_loss: 0.1258 - val_accuracy: 0.9895
```

```
<keras.callbacks.History at 0x7f765a4aa910>
```

Saving the Model

```
model.save('sms_classifier.h5')
```

Preprocessing the Test Dataset

```
test_sequences = tok.texts_to_sequences(X_test)  
test_sequences_matrix = pad_sequences(test_sequences,maxlen=max_len)
```

Testing the Model

```
accr = model.evaluate(test_sequences_matrix,Y_test)
```

```
27/27 [=====] - 1s 24ms/step - loss: 0.1624 -  
accuracy: 0.9868
```

```
print('Test set\n Loss: {:.3f}\n Accuracy:  
{:.3f}'.format(accr[0],accr[1]))
```

Test set

Loss: 0.162

Accuracy: 0.987